

An abstract graphic on the left side of the slide consists of several overlapping circles. The top-left circle is a dark teal color. Below it is a lighter teal circle. To the right of these is a dark grey circle. At the bottom is a large orange circle. The circles overlap in various ways, creating a layered effect.

# A Fast Decision Rule Engine for Anomaly Detection

James Thomas

# Overview

- Introducing a classifier based on one- and two-feature decision rules as an interpretable approach for supervised anomaly detection
  - Practical due to fast implementation
- Pandas API and demo

# Supervised Anomaly Detection Problems

- Binary classification with large class imbalance
  - Normal ML methods can struggle
- Want interpretability because humans often involved in addressing anomaly
  - Why was this classified as an anomaly?

# Decision Rules for Categorical Tabular Data

Cellphone telemetry data:

OS Version	Manufacturer	Device Age	Region	Has error (class)
4.1	Samsung	1	US	1
4.1	Nokia	1	Europe	0
4.2	HTC	3	US	0
4.1	HTC	2	Asia	0
4.1	Nokia	1	Europe	1
4.3	HTC	1	Asia	0

# Decision Rules for Categorical Tabular Data

Cellphone telemetry data:

OS Version	Manufacturer	Device Age	Region	Has error (class)
4.1	Samsung	1	US	1
4.1	Nokia	1	Europe	0
4.2	HTC	3	US	0
4.1	HTC	2	Asia	0
4.1	Nokia	1	Europe	1
4.3	HTC	1	Asia	0

Potential one-feature rule to select anomaly class:

OS Version = 4.1 (Precision 50%, 4 examples)

# Decision Rules for Categorical Tabular Data

Cellphone telemetry data:

OS Version	Manufacturer	Device Age	Region	Has error (class)
4.1	Samsung	1	US	1
4.1	Nokia	1	Europe	0
4.2	HTC	3	US	0
4.1	HTC	2	Asia	0
4.1	Nokia	1	Europe	1
4.3	HTC	1	Asia	0

Potential one-feature rule to select anomaly class:

Manufacturer = Samsung (Precision 100%, 1 example)

# Decision Rules for Categorical Tabular Data

Cellphone telemetry data:

OS Version	Manufacturer	Device Age	Region	Has error (class)
4.1	Samsung	1	US	1
4.1	Nokia	1	Europe	0
4.2	HTC	3	US	0
4.1	HTC	2	Asia	0
4.1	Nokia	1	Europe	1
4.3	HTC	1	Asia	0

Potential two-feature rule to select anomaly class:

OS Version = 4.1 && Device Age = 1 (Precision 66%, 3 examples)

# Decision Rules Are Interpretable

- When limited to one or two features
- Even decision trees (especially deep ones or random forests) and linear models are hard to fully understand



# Extending Decision Rules to Continuous Features

- Find min and max of continuous feature and discretize into equally sized buckets (15 buckets in our system)
- Other discretization schemes possible

# Combining Decision Rules

- Single decision rule unlikely to be enough to classify well
- Create a classifier with many good decision rules; if any of them fires, anomaly is detected (logical OR of rules)
- Still interpretable – human can see which rule(s) fired for particular example

# Evaluating Decision Rules: Counts

- Maintain count of anomalies and total examples for all one-feature and two-feature decision rules

Feature #1	Feature #1 Value	Feature #2	Feature # 2 Value	# Anomalies	# Total Examples	Precision
OS Version =	4.1	--	--	2	4	0.5
OS Version =	4.1	Region =	Asia	0	1	0.0
OS Version =	4.1	Region =	Europe	1	2	0.5

⋮

# Computing Two-Feature Counts

- Gets expensive with large number of features (all pairs)
- We have a fast C++ implementation with experimental GPU/FPGA acceleration available
  - Can scale to the ~1000 feature range for large datasets

# Selecting Decision Rules: Filter on Precision/Count

- Create a classifier with all decision rules having precision  $\geq p\_thresh$  and total examples  $\geq c\_thresh$

# Pruning Decision Rules from Classifier

- Overall classifier will likely have lower than  $p\_thresh$  precision because there will be more overlap in the rules' correctly classified anomalies than in false positives
- Need way to prune redundant rules
  - Fewer rules also easier for human to process

# Pruning Decision Rules from Classifier

- One heuristic: sort selected rules descending by anomaly count
- Iterate through rules and compute incremental precision (new true positives / new false positives) over previous rules
- Discard rules with incremental precision  $< p\_thresh'$  and incremental classified examples  $< c\_thresh'$

# Pruning Decision Rules from Classifier

- With  $p\_thresh'$  very small and  $c\_thresh' = 1$ , eliminates only rules that have no correct incremental classifications
- Strictly improves classifier precision with no change to recall
- Other heuristics possible...



# Pandas API Summary

- `s = compute_sums(train_set, class_name)`
  - Compute all one-feature and two-feature counts
- `r = s.get_rules(p_thresh, c_thresh)`
  - Return dataframe of all rules with precision  $\geq p\_thresh$  and total training examples classified  $\geq c\_thresh$

# Pandas API Summary

- `r = s.prune(r, examples, p_thresh', c_thresh')`
  - Prune rules based on incremental performance; `examples` can just be the training set
- `s.display_rules(r)`
  - Display all rules in human-readable format
- `s.evaluate_summary(r, test_set)`
  - Return precision and recall of classifier consisting of all rules in `r`

# Demo

- [https://github.com/jjthomas/rule\\_engine](https://github.com/jjthomas/rule_engine)
- [jamesjoethomas@gmail.com](mailto:jamesjoethomas@gmail.com)